Haotian Yang (Sora)

Bocconi University, Milan
haotian.yang@studbocconi.it · +39 344 767 9865 · huggingface.co/sora42y · github.com/soray42

EDUCATION

Bocconi University 2025–2028
B.Sc. in Economics, Management and Computer Science (BEMACS), second year.

PUBLICATIONS

Yang, H. 2026. "import pandas as np: the numpy strikes back."
BlackboxNLP Workshop at EMNLP 2026 (accepted; poster presentation, Budapest).
Pretrained aliases beat local bindings: with `import pandas as np`, 74% of completions break vs. 12% under the conventional alias (9 models, 0.5B–8B).

PRESENTATIONS

Yang, H. 2026. "One Shock, Many Voices: Do LLM Managers Discount Correlated Analyst Reports?"
AI4F Workshop at ICAIF '26 (Bocconi, Milan, Nov. 2026). Workshop presentation, non-archival.
An LLM manager applies about a third of the Bayes-correct discount for correlated analyst reports (0.21 vs. 0.68 log-odds); asking for the rationale first raises it to 0.56.

Yang, H., and Y. Li. 2026. "On the recovery of meaningful structure from financial transfer entropy networks." *EcoSta 2026*. Oral presentation (presenting author).

MANUSCRIPTS UNDER REVIEW

Yang, H. 2026. "Do You Copy? When Language Models Overwrite Observed Values with Remembered Defaults."
Under review at *ICLR 2027*.
Shown a well-known service on a non-default port, agents write the remembered default: +14.2 points when the value arrives as a tool result, vs. +1.8 from a retrieved document.

Yang, H., D. Hovy, and G. Glavaš. 2026. "Seductive Details: Decoupling Salience from Relevance in LLM Distractibility."
Short paper, under review at *ACL Rolling Review* (October 2026 cycle).
Making an irrelevant cue more salient, with its claim held fixed, flips the model's verdict on it in 18% of scenario–model pairs, in both directions (4 models, 500 scenarios).

OPEN-SOURCE RELEASES

Sev-2B-preview (Oct 2026). One forward pass, no reasoning tokens: .486 on a Knights & Knaves diagnostic vs. .236 for the best other 2B decision model; .704 on ProverQA. Model, training data, and training/evaluation code (Apache-2.0).

Seductive Details dataset (Oct 2026). Scenarios in 5 languages, model outputs, judge and human labels; one command reproduces every number in the paper.

LEADERSHIP & ACTIVITIES

Black Swan Students Society Jul 2026–present
Founding Member. Conducting academic research in prediction markets.

Bocconi Students Blockchain Association Sep 2025–Jul 2026
Market Analyst.

RESEARCH INTERESTS, SKILLS & LANGUAGES

Research interests: LLM post-training and reasoning distillation, LLM behavior and interpretability, information theory and complex systems, AI in finance.

Technical: Julia, Python, LaTeX, Git, SLURM/HPC.

Languages: Mandarin (native), English (fluent), Japanese (JLPT N3).